

Machine Learning Final Project

Group 1 : 110511109李崇碩、110511277蔡東宏

2024/06/21

Outline

- **Data Augmentation**
- **Model Choosing**
- **ShuffleNet Model**
- **Final Model**
- **Batch Size**
- **Optimizer Choosing**
- **Dynamic adjusting Learning Rate**
- **Reference**

Data Augmentation

■ Data Augmentation is used to avoid overtraining

- Resize (256 x 256)
- RandomHorizontalFlip、RandomVerticalFlip
- RandomRotation
- RandomErasing
- ColorJitter (未採用)
- Normalize



Model Choosing

■ Comparing Different Model (parameters, FLOPs, Testing Accuracy)

Model	Parameter	FLOPs	Testing Accuracy
ResNet18	11.7M	2.382G	0.9266
ResNet50	25.56M	5.4G	0.9266
ShuffleNet_v2	1.367M	57.9M	0.8349
EfficientNet	5.289M	641.64G	0.9083
MobileNet_v2	3.505M	427.346G	0.9174
VGG16	138.36M	20.168G	0.9174
VGG19	143.667M	25.604G	0.9174
DenseNet121	7.98M	3.78G	0.9266

Model Choosing

■ **Without Modifying ShuffleNet_v2**

➤ **Parameters : 1.376M 、 FLOPs : 57.9M 、 Testing Accuracy : 0.8349**

■ **Modifying ShuffleNet_v2**

□ **Removing the latest two layers**

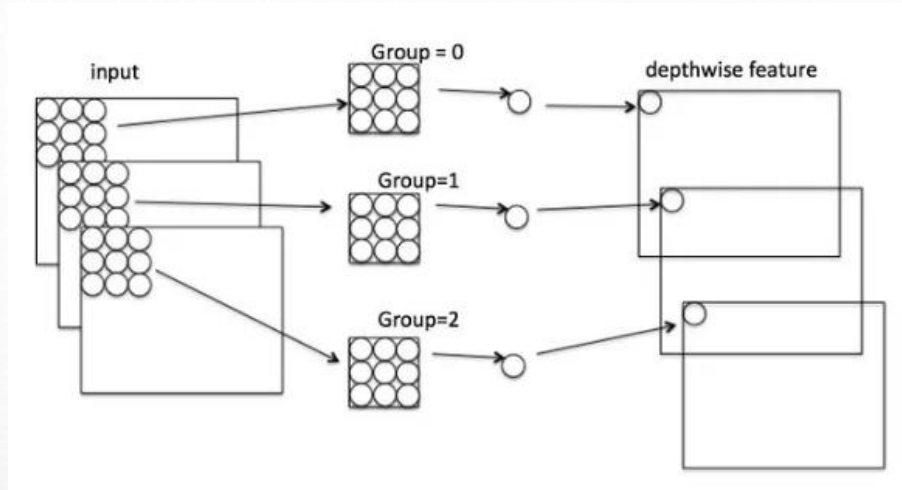

□ **Adding two full connected layers**

□ **Adding some dropout to avoid overtraining**


➤ **Parameters : 0.143M 、 FLOPs : 44M 、 Testing Accuracy : 0.9**

ShuffleNet Model

- Design concept
 - Group Convolution
 - Depthwise and Pointwise Convolution

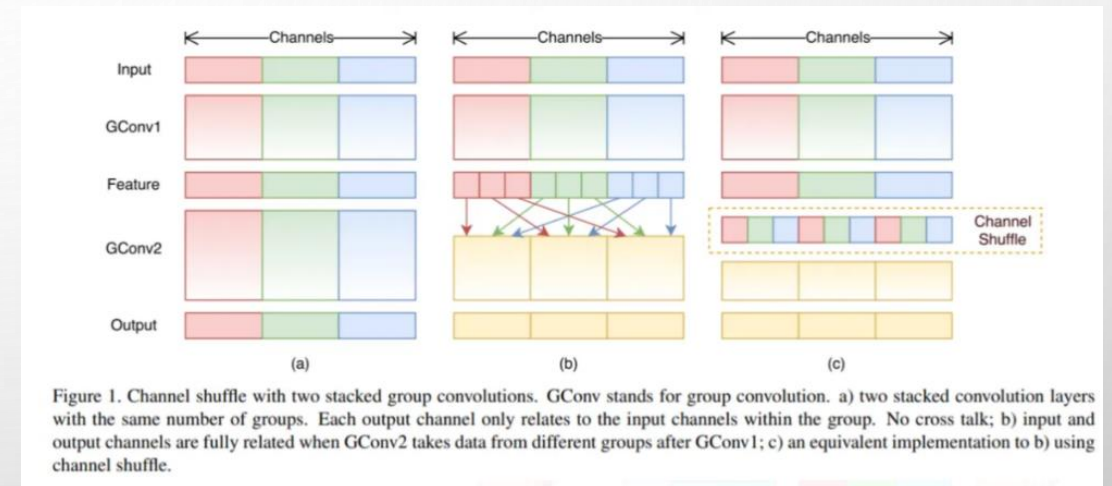


上圖為depthwise feature，下圖為pointwise feature



ShuffleNet V1 Model

- ShuffleNet V1 is design based on the bottleneck of ResNet.
- Using group Convolution and Channel Shuffle to reduce computation and parameter size.
- Channel Shuffle allows information exchange between channels



Channel Shuffle示意圖

ShuffleNet V1 VS ShuffleNet V2

- **Computational efficiency: ShuffleNet_V2 > ShuffleNet_V1.**
- **ShuffleNet_V2 is simpler and more.**
- **ShuffleNet_V2 introduces new computational units (such as SE modules)**
- **ShuffleNet_V2 allocates channels differently based on actual computational characteristics.**

Self-designed Model

nn.Conv2d:

- **Function:** Effectively extracts local features from images. Stacking multiple convolutional layers can progressively extract higher-level features.
- **Design Reason:** After each convolutional layer, the number of channels (features) gradually increases (from 3 (RGB) to 32, 32 to 64, and then to 128). This is because we want to capture more complex features as the layers deepen.

nn.ReLU:

- **Function:** The ReLU function introduces non-linearity, allowing the model to learn more complex features.
- **Design Reason:** Avoid the limitations of a linear model. Additionally, ReLU is computationally efficient and helps solve the vanishing gradient problem.

```
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=2):
        super(SimpleCNN, self).__init__()
        self.features = nn.Sequential([
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ])
        self.classifier = nn.Sequential([
            nn.Dropout(p=0.5),
            nn.Linear(128 * 6 * 6, 200), # Adjusted for .
            nn.ReLU(inplace=True),
            nn.Dropout(p=0.5),
            nn.Linear(200, num_classes)
        ])

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x
```

Self-designed Model

nn.MaxPool2d:

- **Function:** Reduces the size of the feature map by taking the maximum value within a window, preserving the most important features.
- **Design Reason:** Reduce computational load and the number of parameters. Through downsampling, they make the feature map more invariant and increase the receptive field.

nn.Dropout:

- **Function:** Randomly drops some neurons during training.
- **Design Reason:** prevent overfitting and improve the model's generalization ability.

```
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=2):
        super(SimpleCNN, self).__init__()
        self.features = nn.Sequential([
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2)
        ])
        self.classifier = nn.Sequential(
            nn.Dropout(p=0.5),
            nn.Linear(128 * 6 * 6, 200), # Adjusted for .
            nn.ReLU(inplace=True),
            nn.Dropout(p=0.5),
            nn.Linear(200, num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x
```

Self-designed Model

■ **Result :** FLOPs: 24146320.0, Params: 1015450.0

■ **Conclusion :** found that even with a simple method without stacking many layers, the parameters was already suboptimal. Therefore, We decided not to continue testing it and instead replaced it with ShuffleNet.

Epoch 53/100, Loss: 0.2989, Accuracy: 0.8786, Test Loss: 0.6106, Test Accuracy: 0.7167
Epoch 54/100, Loss: 0.2693, Accuracy: 0.8946, Test Loss: 0.6460, Test Accuracy: 0.6917
Epoch 55/100, Loss: 0.2715, Accuracy: 0.8911, Test Loss: 0.6245, Test Accuracy: 0.6917
Epoch 56/100, Loss: 0.2708, Accuracy: 0.8911, Test Loss: 0.6230, Test Accuracy: 0.7250
Epoch 57/100, Loss: 0.2830, Accuracy: 0.8804, Test Loss: 0.6079, Test Accuracy: 0.7167
Epoch 58/100, Loss: 0.2778, Accuracy: 0.8821, Test Loss: 0.6180, Test Accuracy: 0.6833
Epoch 59/100, Loss: 0.2763, Accuracy: 0.9000, Test Loss: 0.6277, Test Accuracy: 0.7000
Epoch 60/100, Loss: 0.2850, Accuracy: 0.9036, Test Loss: 0.6211, Test Accuracy: 0.7083
Epoch 61/100, Loss: 0.2715, Accuracy: 0.8946, Test Loss: 0.6094, Test Accuracy: 0.6917
Epoch 62/100, Loss: 0.2768, Accuracy: 0.8946, Test Loss: 0.5993, Test Accuracy: 0.7083
Epoch 63/100, Loss: 0.2744, Accuracy: 0.8982, Test Loss: 0.6210, Test Accuracy: 0.7000
Epoch 64/100, Loss: 0.2795, Accuracy: 0.8857, Test Loss: 0.6162, Test Accuracy: 0.7083
Epoch 65/100, Loss: 0.2695, Accuracy: 0.8946, Test Loss: 0.6280, Test Accuracy: 0.7000
Epoch 66/100, Loss: 0.2693, Accuracy: 0.9036, Test Loss: 0.6373, Test Accuracy: 0.6917
Epoch 67/100, Loss: 0.2821, Accuracy: 0.8911, Test Loss: 0.6073, Test Accuracy: 0.7000
Epoch 68/100, Loss: 0.2933, Accuracy: 0.8821, Test Loss: 0.6103, Test Accuracy: 0.7083
Epoch 69/100, Loss: 0.2852, Accuracy: 0.8839, Test Loss: 0.6318, Test Accuracy: 0.7083
Epoch 70/100, Loss: 0.2856, Accuracy: 0.8768, Test Loss: 0.6266, Test Accuracy: 0.7000
Epoch 71/100, Loss: 0.2726, Accuracy: 0.8911, Test Loss: 0.6260, Test Accuracy: 0.7000
Epoch 72/100, Loss: 0.2694, Accuracy: 0.8893, Test Loss: 0.6320, Test Accuracy: 0.7000
Epoch 73/100, Loss: 0.2616, Accuracy: 0.9000, Test Loss: 0.6227, Test Accuracy: 0.7167
Epoch 74/100, Loss: 0.2653, Accuracy: 0.8982, Test Loss: 0.6032, Test Accuracy: 0.7167
Epoch 75/100, Loss: 0.2762, Accuracy: 0.8929, Test Loss: 0.6217, Test Accuracy: 0.7000
Epoch 76/100, Loss: 0.2660, Accuracy: 0.9018, Test Loss: 0.5976, Test Accuracy: 0.7333
Epoch 77/100, Loss: 0.2811, Accuracy: 0.8982, Test Loss: 0.6120, Test Accuracy: 0.7083
Epoch 78/100, Loss: 0.2861, Accuracy: 0.9089, Test Loss: 0.6007, Test Accuracy: 0.7167
Epoch 79/100, Loss: 0.2730, Accuracy: 0.9000, Test Loss: 0.6171, Test Accuracy: 0.7167
Epoch 80/100, Loss: 0.2798, Accuracy: 0.8946, Test Loss: 0.6179, Test Accuracy: 0.7167
Epoch 81/100, Loss: 0.2676, Accuracy: 0.9000, Test Loss: 0.5936, Test Accuracy: 0.7083
Epoch 82/100, Loss: 0.2839, Accuracy: 0.8821, Test Loss: 0.6056, Test Accuracy: 0.7083
Epoch 83/100, Loss: 0.2705, Accuracy: 0.8964, Test Loss: 0.6238, Test Accuracy: 0.7250
Epoch 84/100, Loss: 0.2702, Accuracy: 0.8946, Test Loss: 0.5975, Test Accuracy: 0.7083
Epoch 85/100, Loss: 0.2748, Accuracy: 0.9036, Test Loss: 0.5957, Test Accuracy: 0.7000
Epoch 86/100, Loss: 0.2780, Accuracy: 0.8982, Test Loss: 0.6026, Test Accuracy: 0.7167
Epoch 87/100, Loss: 0.2757, Accuracy: 0.8964, Test Loss: 0.6325, Test Accuracy: 0.7000
Epoch 88/100, Loss: 0.2872, Accuracy: 0.8911, Test Loss: 0.6124, Test Accuracy: 0.7250
Epoch 89/100, Loss: 0.2577, Accuracy: 0.9107, Test Loss: 0.6264, Test Accuracy: 0.6917
Epoch 90/100, Loss: 0.2671, Accuracy: 0.8893, Test Loss: 0.6187, Test Accuracy: 0.7083
Epoch 91/100, Loss: 0.2709, Accuracy: 0.8982, Test Loss: 0.6158, Test Accuracy: 0.7167
Epoch 92/100, Loss: 0.2562, Accuracy: 0.8929, Test Loss: 0.6230, Test Accuracy: 0.7083
Epoch 93/100, Loss: 0.2725, Accuracy: 0.9071, Test Loss: 0.6152, Test Accuracy: 0.7083
Epoch 94/100, Loss: 0.2688, Accuracy: 0.8964, Test Loss: 0.6153, Test Accuracy: 0.7167
Epoch 95/100, Loss: 0.2740, Accuracy: 0.9000, Test Loss: 0.6295, Test Accuracy: 0.6917
Epoch 96/100, Loss: 0.2821, Accuracy: 0.8982, Test Loss: 0.6078, Test Accuracy: 0.7250
Epoch 97/100, Loss: 0.2834, Accuracy: 0.8804, Test Loss: 0.6199, Test Accuracy: 0.7083
Epoch 98/100, Loss: 0.2771, Accuracy: 0.8875, Test Loss: 0.6018, Test Accuracy: 0.7250

Final Model

- Removing the latest two layer of ShuffleNet_v2 Model
- Adding global average pool to reduce the parameters of model
- Adding a dropout layer with a probability of 0.2 to randomly drop some neurons

```
class Network(nn.Module):
    def __init__(self, num_classes=2):
        super(Network, self).__init__()
        self.model = torchvision.models.shufflenet_v2_x0_5(pretrained=True)

        self.model = nn.Sequential(*list(self.model.children())[:-2])
        self.model.add_module('global_avg_pool', nn.AdaptiveAvgPool2d(1))
        self.model.add_module('flatten', nn.Flatten())
        self.model.add_module('dropout', nn.Dropout(p=0.2))
        self.model.add_module('fc', nn.Linear(in_features=192, out_features=2))

        self.dropout = nn.Dropout(p=0.02)
        self.fc = nn.Linear(in_features=2, out_features=num_classes)

    def forward(self, x):
        x = self.model(x)
        x = self.dropout(x)
        x = self.fc(x)
        return x
```


Final Model

- Adding a full connected layer
(in_features = 192, out_features = 2)
- Adding a dropout layer with a probability of 0.02 to randomly drop some neurons
- Adding a full connected layer
(in_features = 2, out_features = 2)

```
class Network(nn.Module):  
    def __init__(self, num_classes=2):  
        super(Network, self).__init__()  
        self.model = torchvision.models.shufflenet_v2_x0_5(pretrained=True)  
  
        self.model = nn.Sequential(*list(self.model.children())[:-2])  
        self.model.add_module('global_avg_pool', nn.AdaptiveAvgPool2d(1))  
        self.model.add_module('flatten', nn.Flatten())  
        self.model.add_module('dropout', nn.Dropout(p=0.2))  
        self.model.add_module('fc', nn.Linear(in_features=192, out_features=2))  
  
        self.dropout = nn.Dropout(p=0.02)  
        self.fc = nn.Linear(in_features=2, out_features=num_classes)  
  
    def forward(self, x):  
        x = self.model(x)  
        x = self.dropout(x)  
        x = self.fc(x)  
        return x
```

Calculate Performance

- Use the profile function from thop package to calculate FLOPs and parameters.
- Additionally, use the summary function from torchsummary package to print out the parameter count, **layer type**, and **output shape** of each layer.

```
from thop import profile
from torchsummary import summary
model_for_profiling = copy.deepcopy(model)
# Profile the cloned model without modifying the original model's state_dict
model_for_profiling.eval() # Set the cloned model to evaluation mode
with torch.no_grad():
    flops, params = profile(model_for_profiling, inputs=(random_tensor,), verbose=False)
print(f"FLOPs: {flops}, Params: {params}")
summary(model, (3, 256, 256))
```

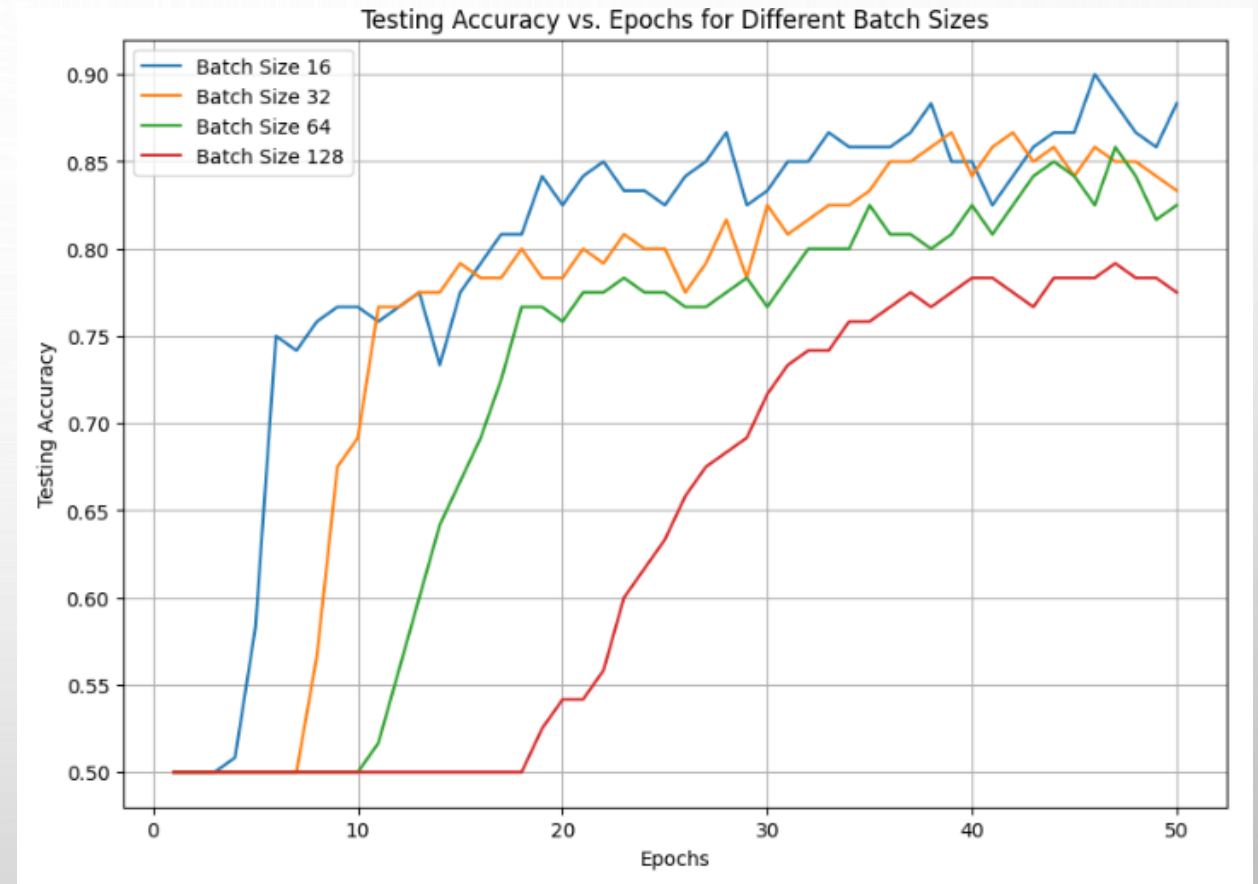

Seed

- To ensure that the testing environment remains consistent
- results before and after code optimizations can achieve accurate comparison

```
def same_seeds(seed):  
    random.seed(seed)  
    np.random.seed(seed)  
    torch.manual_seed(seed)  
    if torch.cuda.is_available():  
        torch.cuda.manual_seed(seed)  
        torch.cuda.manual_seed_all(seed)  
    torch.backends.cudnn.benchmark = False  
    torch.backends.cudnn.deterministic = True  
  
same_seeds(48763)
```

Batch Size

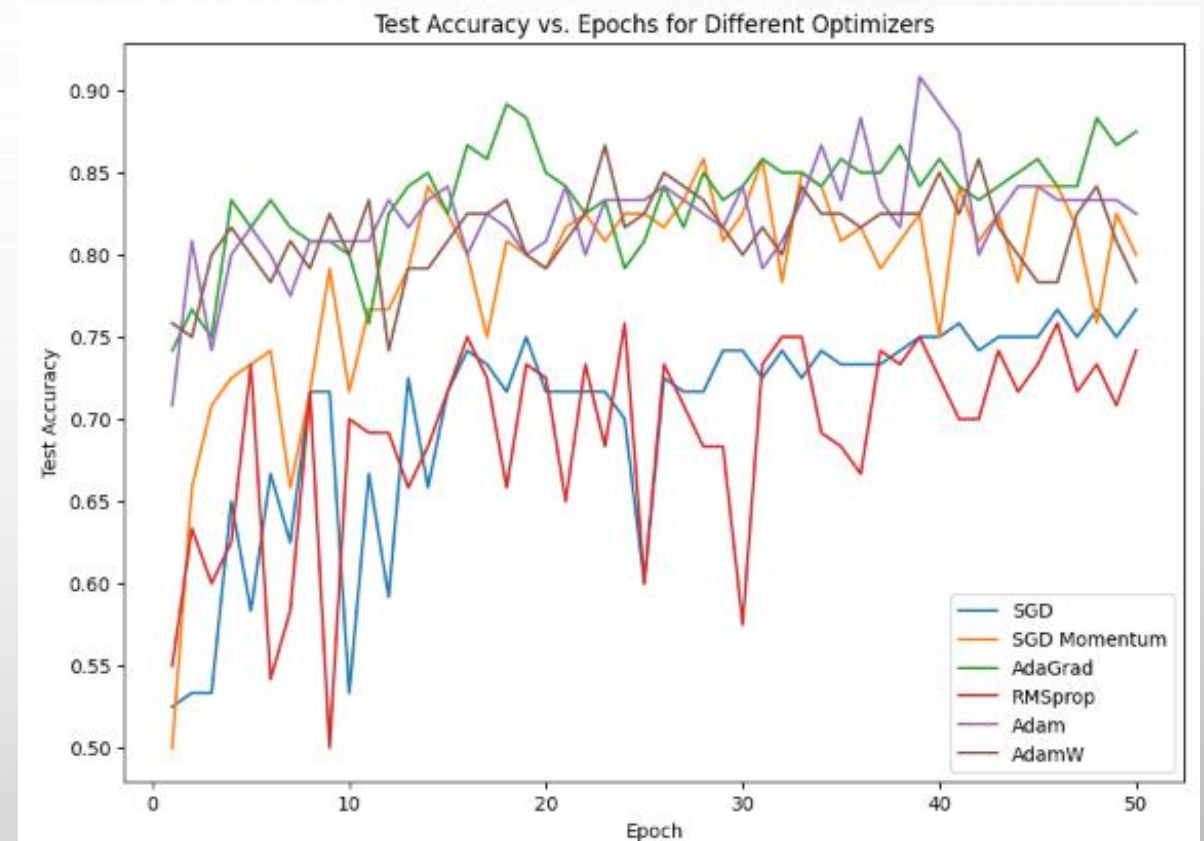
- **Comparing Different Batch size**
 - **A smaller batch size converges faster and results in higher accuracy.**
- **We finally using 16 as our batch size**



Optimizer Choosing


■ Comparing Different Optimizer

- SGD
- SGD Momentum
- AdaGrad
- RMSprop
- Adam
- AdamW




Optimizer Choosing

- Through many times testing ,
We find that set Learning Rate
= 0.0001 can achieve the best
accuracy
- Although AdaGrad surpassed
Adam in the final few epochs,
Adam achieved a maximum
testing accuracy of 90%.
Therefore, we chose Adam as
the optimizer.



Dynamic adjusting Learning Rate

- Using ReduceLROnPlateau to dynamic adjusting learning rate.
- We can observe that in first 40 epochs, the learning rate is 0.0001. However, in final few epochs, the learning rate is reduced.



Dynamic adjusting Learning Rate

- After lowering the learning rate, the large-scale instability and oscillations were significantly reduced.

Before Adjusting

```
Epoch 40/50, Loss: 0.1276, Accuracy_train: 0.9625, Accuracy_test: 0.8417,  
Epoch 41/50, Loss: 0.1223, Accuracy_train: 0.9625, Accuracy_test: 0.8417,  
Epoch 42/50, Loss: 0.1410, Accuracy_train: 0.9589, Accuracy_test: 0.8417,  
Epoch 43/50, Loss: 0.1193, Accuracy_train: 0.9661, Accuracy_test: 0.8250,  
Epoch 44/50, Loss: 0.1028, Accuracy_train: 0.9750, Accuracy_test: 0.8333,  
Epoch 45/50, Loss: 0.1113, Accuracy_train: 0.9589, Accuracy_test: 0.8250,  
Epoch 46/50, Loss: 0.0831, Accuracy_train: 0.9786, Accuracy_test: 0.8417,  
Epoch 47/50, Loss: 0.0933, Accuracy_train: 0.9661, Accuracy_test: 0.8250,  
Epoch 48/50, Loss: 0.0842, Accuracy_train: 0.9786, Accuracy_test: 0.8500,  
Epoch 49/50, Loss: 0.0785, Accuracy_train: 0.9732, Accuracy_test: 0.8250,  
Epoch 50/50, Loss: 0.0778, Accuracy_train: 0.9821, Accuracy_test: 0.8333,
```

After Adjusting

```
Epoch 00041: reducing learning rate of group 0 to 1.0000e-05.  
Learning Rate: 0.000010  
Epoch 42/50, Loss: 0.2302, Accuracy_train: 0.9393, Accuracy_test: 0.8500  
Learning Rate: 0.000010  
Epoch 43/50, Loss: 0.2496, Accuracy_train: 0.9107, Accuracy_test: 0.8583  
Learning Rate: 0.000010  
Epoch 44/50, Loss: 0.2406, Accuracy_train: 0.9286, Accuracy_test: 0.8500  
Learning Rate: 0.000010  
Epoch 45/50, Loss: 0.2355, Accuracy_train: 0.9304, Accuracy_test: 0.8583  
Learning Rate: 0.000010  
Epoch 46/50, Loss: 0.2483, Accuracy_train: 0.9214, Accuracy_test: 0.8583  
Learning Rate: 0.000010  
Epoch 47/50, Loss: 0.2367, Accuracy_train: 0.9268, Accuracy_test: 0.8583  
Learning Rate: 0.000010  
Epoch 48/50, Loss: 0.2319, Accuracy_train: 0.9321, Accuracy_test: 0.8500  
Learning Rate: 0.000010  
Epoch 49/50, Loss: 0.2348, Accuracy_train: 0.9321, Accuracy_test: 0.8583  
Epoch 00049: reducing learning rate of group 0 to 1.0000e-06.  
Learning Rate: 0.000001  
Epoch 50/50, Loss: 0.2224, Accuracy_train: 0.9304, Accuracy_test: 0.8583  
Training complete
```


Recording Best Accuracy Model

- Record the Testing Accuracy for each run, then load the model with the best Testing Accuracy for the final demo to ensure optimal accuracy.

```
model = Network(num_classes=2)
model.load_state_dict(torch.load('best_model.pth'))
model = model.to(device)
model.eval()

correct_test = 0
total_test = 0
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total_test += labels.size(0)
        correct_test += (predicted == labels).sum().item()

epoch_acc_test = correct_test / total_test
print(f'Accuracy_test: {epoch_acc_test:.4f}')
```

Reference

- <https://hackmd.io/@allen108108/H1l4zqtp4>
- <https://medium.com/%E5%AD%B8%E4%BB%A5%E5%BB%A3%E6%89%8D/note-shufflenet-%E5%AD%B8%E7%BF%92%E5%BF%83%E5%BE%97-38f2f715e19a>